

RETAILMART V3 • SQL

DDL & DML

WhatsApp Announcement

DDL & DML: Build Your First Tables, Then Fill Them!

Yesterday you explored RetailMart V3 - 55 tables, millions of rows. You looked at the data. Today you go from spectator to architect AND builder: first you BUILD tables (DDL), then you put DATA into them (DML).

Today's Focus:

- CREATE DATABASE, CREATE SCHEMA, CREATE TABLE - building structure from scratch
- Every column is TEXT today - we focus 100% on STRUCTURE, not types (types come tomorrow)
- DROP, IF NOT EXISTS, naming conventions
- INSERT - add rows. UPDATE - change rows. DELETE - remove rows. TRUNCATE - empty a table

Important: ALL building and data-entry today happens in YOUR practice database (accio_NN), NOT in RetailMart. You never change RetailMart's tables - you observe its design and build your own.

Course Portal: <https://starlordali4444.github.io/postgresql/>

Today you write your FIRST real SQL and put your FIRST rows in. Let's go!

- Siraj Sir

#Day3 #DDL #DML #PostgreSQL

TODAY'S AGENDA (2 Hours)

Section	Time	Topics
Quick Recap	5 mins	5 SQL categories, Database hierarchy, pgAdmin connection check
Topic 1: DDL - Building Structures	35 mins	CREATE DATABASE / SCHEMA / TABLE, Why every column is TEXT today, DROP, IF NOT EXISTS, Naming conventions
Practice Block 1	20 mins	4 problems: Build tables from scratch (all TEXT columns)
Topic 2: DML - Putting Data In	40 mins	INSERT (single, multi-row, RETURNING), UPDATE (always WHERE!), DELETE, TRUNCATE, DROP vs TRUNCATE vs DELETE
Practice Block 2	25 mins	4 problems: INSERT / UPDATE / DELETE / TRUNCATE on your tables
Wrap-up	15 mins	Summary, self-check, homework, Data Types & Constraints preview

5-Minute Memory Check

From SQL Intro:

- 5 SQL Categories: DDL, DML, DQL, DCL, TCL
- Database Hierarchy: Server -> Database -> Schema -> Table -> Row -> Cell
- SQL is DECLARATIVE - you say WHAT, engine decides HOW

From Setup & Onboarding:

- PostgreSQL = server process. pgAdmin = GUI client. psql = terminal client.
- RetailMart V3 loaded - 16 schemas, 55 tables
- You ran `SELECT count(*)` and `SELECT * ... LIMIT 5`
- You browsed schemas: customers, sales, products, stores, finance...

Today's Big Question:

'Those 55 tables did not appear by magic.
Someone wrote `CREATE TABLE` 55 times (DDL).
Someone wrote millions of `INSERTs` to fill them (DML).
Today, YOU do both - on your OWN tables.'

The Architect, the Builder, and the Resident (1/2)

Three people work on a building:

The Architect (DDL):

Draws the blueprint. How many floors? How many rooms? What is each room called? The architect defines STRUCTURE - she never moves furniture in.

The Builder (DML):

Follows the blueprint and puts things INSIDE - moves furniture in (INSERT), rearranges it (UPDATE), removes it (DELETE). The builder cannot add a new room - that needs the architect.

The Architect, the Builder, and the Resident (2/2)

The Resident (DQL):

Lives in the house and looks around (SELECT). Changes nothing. You meet the Resident starting tomorrow-ish (Filtering & Sorting).

Today you are BOTH the Architect (Topic 1: DDL) and the Builder (Topic 2: DML).

[Mapping to SQL]

Architect draws blueprint	-> CREATE TABLE	(define structure)	[Topic 1]
Architect modifies blueprint	-> ALTER TABLE	(change structure)	[tomorrow]
Architect demolishes a room	-> DROP TABLE	(destroy structure)	[Topic 1]
Builder brings furniture	-> INSERT	(add data)	[Topic 2]
Builder rearranges furniture	-> UPDATE	(change data)	[Topic 2]
Builder removes furniture	-> DELETE	(remove data)	[Topic 2]
Resident looks around	-> SELECT	(read data)	[later]

Today's Big Idea: Every Column Is TEXT

A table is a STRUCTURE: named columns, in order, grouped into a schema. Every column must have a TYPE - and today we give EVERY column the simplest, most forgiving type: TEXT. TEXT accepts any characters - letters, digits, dates typed as text, anything.

Why TEXT for everything today?

- So you focus 100% on STRUCTURE - what tables exist, what they are named, what columns they hold, how schemas organize them.
- Choosing the RIGHT type for each column (INT, NUMERIC, DATE, BOOLEAN...) is a whole skill - that is TOMORROW (Data Types).
- Tomorrow you will ALTER these TEXT columns into their proper types - motivated by the real data you INSERT today.

So today's CREATE TABLE is beautifully simple: 'column_name TEXT' for every column. No types to agonize over, no rules yet. Just structure.

CREATE DATABASE - The Empty Plot of Land

A database is the top-level container. Everything lives inside it - schemas, tables, views, indexes, functions.

[Syntax & Examples]

-- Basic syntax

```
CREATE DATABASE database_name;
```

-- Create your practice database

```
CREATE DATABASE accio_NN;
```

-- What just happened under the hood?

-- 1. PostgreSQL created a NEW directory on disk

-- 2. Copied the template database (template1) as the skeleton

-- 3. Created one default schema inside: 'public'

-- 4. Created system catalogs (pg_catalog, information_schema)

-- 5. The database has ZERO user tables. You build those.

-- Connect to it:

```
\c accio_NN      -- in psql
```

-- or: right-click in pgAdmin -> Query Tool

CREATE SCHEMA - Folders Inside the Database (1/4)

A schema is a namespace - a logical folder that groups related tables. Without schemas, all your tables dump into one flat 'public' schema. Imagine a folder with 55 files and no subfolders - chaos.

[Syntax & Examples]

```
CREATE SCHEMA schema_name;
```

```
-- Organize like a real company
```

```
CREATE SCHEMA shop;           -- products, inventory, suppliers  
CREATE SCHEMA people;        -- customers, employees  
CREATE SCHEMA transactions;   -- orders, payments, returns
```

CREATE SCHEMA - Folders Inside the Database (2/4)

```
-- Why schemas matter:  
-- 1. Organization: many tables in a few folders vs all loose  
-- 2. Namespace:    shop.products is different from analytics.products  
-- 3. Clarity:       SELECT * FROM sales.orders - you KNOW where it lives
```

CREATE SCHEMA - Folders Inside the Database (3/4)

[RetailMart V3's 16 Schemas - Why Each Exists]

Schema	Tables	Purpose
core	6	Static dimensions (dates, regions, categories)
products	4	What we sell (items, suppliers, inventory)
stores	3	Where we sell (locations, staff, expenses)
customers	5	Who buys (profiles, addresses, reviews, wallets)
sales	5	The transactions (orders, items, payments, shipments)
finance	6	Money tracking (accounts, transfers, revenue)
hr	2	People ops (attendance, salary history)
marketing	3	Campaign performance (ads, email clicks)
support	1	Customer complaints (tickets)
web_events	2	Digital tracking (page views, events)
supply_chain	3	Logistics (warehouses, shipments, stock)
loyalty	3	Rewards program (tiers, members, redemptions)
manufacture	2	Production (lines, work orders)
payroll	2	Salary processing (tax brackets, pay slips)
call_center	2	Phone support (calls, transcripts)
audit	6	System logs (API requests, changes, errors)

CREATE SCHEMA - Folders Inside the Database (4/4)

Total: 16 schemas, 55 tables. Every table has a HOME.

CREATE TABLE - The Main Event (1/5)

This is the DDL command you will use most. It defines a table - its name, its columns, and the data type of each column. Every column **MUST** have a type. Today, that type is always TEXT.

[Syntax - The Template]

```
CREATE TABLE schema_name.table_name (  
    column_1 TEXT,  
    column_2 TEXT,  
    column_3 TEXT,  
    ...  
);
```

```
-- If you skip schema_name, the table goes into the 'public' schema.  
-- ALWAYS specify the schema explicitly. It is a professional habit.
```

CREATE TABLE - The Main Event (2/5)

[Example 1 - A Students Table (all TEXT)]

```
CREATE TABLE public.students (  
    student_id    TEXT,  
    first_name    TEXT,  
    last_name     TEXT,  
    email         TEXT,  
    date_of_birth TEXT,  
    is_active     TEXT  
);
```

CREATE TABLE - The Main Event (3/5)

```
-- What happened?  
-- PostgreSQL created a table with 6 columns, all TEXT.  
-- Each column has a NAME and a TYPE (TEXT for now).  
-- The table has 0 rows. It is an empty blueprint waiting for data.  
-- Tomorrow: student_id -> INT, date_of_birth -> DATE, is_active -> BOOLEAN.
```

CREATE TABLE - The Main Event (4/5)

[Example 2 - A Products Table (inspired by RetailMart)]

```
CREATE TABLE shop.products (  
    product_id    TEXT,      -- tomorrow: INT (auto-ID with SERIAL)  
    product_name  TEXT,  
    brand_name    TEXT,  
    price         TEXT,      -- tomorrow: NUMERIC(10,2) for money  
    cost_price    TEXT,      -- tomorrow: NUMERIC(10,2)  
    category      TEXT,  
    is_available  TEXT       -- tomorrow: BOOLEAN  
);
```

```
-- Notice every column is TEXT. That is intentional.  
-- We are practising STRUCTURE today, not type selection.
```


CREATE TABLE - The Main Event (5/5)

[Example 3 - An Orders Table (core of any e-commerce DB)]

```
CREATE TABLE transactions.orders (  
  order_id      TEXT,  
  customer_id   TEXT,    -- will LINK to customers tomorrow (FOREIGN KEY)  
  store_id      TEXT,  
  order_date    TEXT,    -- tomorrow: DATE  
  order_status  TEXT,  
  gross_total   TEXT,    -- tomorrow: NUMERIC(12,2)  
  net_total     TEXT,  
  payment_mode  TEXT  
);
```

-- customer_id and store_id will REFERENCE other tables.

-- We make those links real (FOREIGN KEY) tomorrow. For now: plain TEXT.

Naming Conventions - The Unwritten Rules

[Professional Naming Standards]

RULE	Good	Bad
Use snake_case	first_name	firstName, FirstName
Use lowercase only	order_date	Order_Date, ORDER_DATE
Table = plural noun	customers, products	customer, product_tbl
Column = singular	city, order_status	cities, statuses
ID column = table_id	customer_id, product_id	id, cust, pid
Boolean = is_ / has_ prefix	is_active, has_projector	active, projector
Timestamp = _at suffix	created_at, updated_at	creation, update_time
Date = _date or _on suffix	order_date, enrolled_on	date, when
No reserved words	order_status	order, status, type

```
-- Why snake_case?  
-- PostgreSQL internally lowercases everything.  
-- 'OrderDate' becomes 'orderdate' unless you quote it. 'order_date' stays as-is.
```

```
-- RetailMart V3 follows these conventions exactly:  
-- customer_id, first_name, order_date, is_default.
```

CREATE TABLE IF NOT EXISTS - The Safety Net

[Why You Need This]

-- Without IF NOT EXISTS:

```
CREATE TABLE shop.products (product_id TEXT, product_name TEXT);
```

-- Run it once -> SUCCESS

-- Run it again -> ERROR: relation 'shop.products' already exists

-- With IF NOT EXISTS:

```
CREATE TABLE IF NOT EXISTS shop.products (  
    product_id    TEXT,  
    product_name TEXT  
);
```

-- Run it once -> SUCCESS

-- Run it again -> NOTICE: relation already exists, skipping

-- No error. Safe to run 100 times.

-- ALWAYS use this in scripts that might be run repeatedly:

-- setup scripts, migration scripts, CI/CD pipelines need idempotency.

DROP - The Demolition Crew

DROP permanently destroys a database object. There is no undo. No recycle bin. No Ctrl+Z. Gone forever.

[The DROP Commands]

-- Drop a table

```
DROP TABLE shop.products;
```

-- Drop a schema (fails if schema still has tables)

```
DROP SCHEMA shop;
```

-- ERROR: cannot drop schema shop because other objects depend on it

-- Drop schema AND everything inside it

```
DROP SCHEMA shop CASCADE;
```

-- CASCADE = 'destroy the folder AND all tables inside'. DANGEROUS.

-- Drop a database (must be connected to a DIFFERENT database)

```
DROP DATABASE accio_NN;
```

-- Safe versions (no error if it does not exist)

```
DROP TABLE IF EXISTS shop.products;
```

```
DROP SCHEMA IF EXISTS shop CASCADE;
```

DDL vs DML - The Line That Confuses Everyone

Question Students Ask	Answer
'DROP removes data, right?'	DROP removes the TABLE ITSELF + data. DELETE removes ROWS but the table stays.
'TRUNCATE is DDL or DML?'	TRUNCATE is DDL. It removes ALL rows but keeps the table structure intact.
DROP TABLE	-> Table is GONE. Columns gone. Data gone. Everything gone.
TRUNCATE	-> Table EXISTS but is EMPTY. Structure intact. 0 rows.
DELETE FROM	-> Table exists. Some rows removed. Others remain.
DROP	= Demolish the room
TRUNCATE	= Remove ALL furniture from the room (room stays)
DELETE	= Remove SPECIFIC furniture (room + other furniture stays)

PRACTICE BLOCK 1 - DDL: Building Structures

All practice in YOUR database (accio_NN). Every column is TEXT today - focus on structure. Do NOT touch RetailMart.

SCENARIO: You are the new Database Administrator at Prestige Engineering College. The dean wants to digitize all paper records. Start by building the structure from scratch.

YOUR TASK: Run the following DDL in `accio_NN`:

a) : Create a schema called college.

b) : Create college.students with columns: student_id, first_name, last_name, email, date_of_birth, is_active. All columns TEXT.

c) : Create college.courses with: course_id, course_name, credits, department. All TEXT.

d) : Verify both tables exist: `SELECT table_schema, table_name FROM information_schema.tables WHERE table_schema = 'college';`

e) : Inspect the columns: `\d college.students` (in psql) or pgAdmin's Columns tab.

Hint: Use `CREATE TABLE IF NOT EXISTS` for safety. Mistake? `DROP TABLE IF EXISTS` and recreate. That is the point of a practice database.

Answer



ANSWER

```
CREATE SCHEMA IF NOT EXISTS college;
```

```
CREATE TABLE IF NOT EXISTS college.students (  
    student_id    TEXT,  
    first_name    TEXT,  
    last_name     TEXT,  
    email         TEXT,  
    date_of_birth TEXT,  
    is_active     TEXT  
);
```

```
CREATE TABLE IF NOT EXISTS college.courses (  
    course_id    TEXT,  
    course_name  TEXT,  
    credits      TEXT,  
    department   TEXT  
);
```

```
SELECT table_schema, table_name  
FROM information_schema.tables  
WHERE table_schema = 'college';
```


SCENARIO: The college system is growing. The dean now wants to track teachers, classrooms, and fee payments. Design the tables a real college would need.

YOUR TASK:

a) : Create college.teachers with: teacher_id, first_name, last_name, email, phone, subject_specialty, monthly_salary, joining_date, is_hod. All TEXT.

b) : Create college.classrooms with: room_id, building_name, room_number, capacity, has_projector, has_ac. All TEXT.

c) : Create college.fee_payments with: payment_id, student_id, fee_type, amount, payment_date, payment_mode, academic_year, semester. All TEXT.

d) : List all tables in the college schema. How many do you have now?

Hint: Right now monthly_salary, amount, capacity etc. are all TEXT. Tomorrow you will ALTER them to NUMERIC / INT once real data shows what they need to be.

Answer (1/2)

✓ ANSWER

```
CREATE TABLE IF NOT EXISTS college.teachers (  
    teacher_id      TEXT,  
    first_name      TEXT,  
    last_name       TEXT,  
    email           TEXT,  
    phone           TEXT,  
    subject_specialty TEXT,  
    monthly_salary  TEXT,  
    joining_date    TEXT,  
    is_hod          TEXT  
);
```

```
CREATE TABLE IF NOT EXISTS college.classrooms (  
    room_id      TEXT,  
    building_name TEXT,  
    room_number  TEXT,  
    capacity     TEXT,  
    has_projector TEXT,  
    has_ac       TEXT  
);
```

```
CREATE TABLE IF NOT EXISTS college.fee_payments (  
    fee_id      TEXT,  
    student_id  TEXT,  
    amount      TEXT,  
    payment_date TEXT,  
    is_paid     TEXT  
);
```

Answer (2/2)

✓ ANSWER

```
payment_id    TEXT,  
student_id    TEXT,  
fee_type      TEXT,  
amount        TEXT,  
payment_date  TEXT,  
payment_mode  TEXT,  
academic_year TEXT,  
semester      TEXT  
);  
  
SELECT count(*) FROM information_schema.tables  
WHERE table_schema = 'college'; -- 5 tables now
```

SCENARIO: The Operations Manager at RetailMart says: 'Our supply chain team needs a new tracking system. Build me the schema in YOUR practice database.' You get only the business requirement - no column list.

YOUR TASK:

- a)** : Create a schema called supply in accio_NN.
- b)** : Design supply.warehouses - think: id, name, city, capacity, manager_name, is_active, established_date. All TEXT.
- c)** : Design supply.suppliers - company_name, contact_person, email, phone, city, state, registration_date. All TEXT.
- d)** : Design supply.shipments - shipment_id, supplier_id, warehouse_id, product_count, shipped_date, expected_arrival_date, actual_arrival_date, status. All TEXT.
- e)** : Run \d supply.shipments. Then compare to RetailMart's supply_chain.shipments (\d supply_chain.shipments). What did you miss? What did you add?

Hint: There are no wrong answers in STRUCTURE - only trade-offs. The point is to decide which COLUMNS each table needs. Types come tomorrow.

Answer (1/2)



ANSWER

```
CREATE SCHEMA IF NOT EXISTS supply;

CREATE TABLE IF NOT EXISTS supply.warehouses (
    warehouse_id    TEXT,
    name             TEXT,
    city             TEXT,
    capacity         TEXT,
    manager_name     TEXT,
    is_active        TEXT,
    established_date TEXT
);

CREATE TABLE IF NOT EXISTS supply.suppliers (
    supplier_id      TEXT,
    company_name      TEXT,
    contact_person    TEXT,
    email             TEXT,
    phone             TEXT,
    city              TEXT,
    state             TEXT,
    registration_date TEXT
);
```

Answer (2/2)

✓ ANSWER

```
CREATE TABLE IF NOT EXISTS supply.shipments (  
  shipment_id      TEXT,  
  supplier_id      TEXT,    -- will link -> suppliers tomorrow  
  warehouse_id     TEXT,    -- will link -> warehouses tomorrow  
  product_count    TEXT,  
  shipped_date      TEXT,  
  expected_arrival_date TEXT,  
  actual_arrival_date TEXT,  
  status           TEXT  
);
```

SCENARIO: The CTO says: 'We are prototyping RetailMart V4. Recreate the sales schema from scratch - just by looking at the current V3 design. Prove you understand DDL and naming.'

YOUR TASK:

- a)** : Create a schema called mini_mart in accio_NN.
- b)** : Look at RetailMart's sales.orders (\d sales.orders). Recreate it as mini_mart.orders with the SAME column names. All columns TEXT.
- c)** : Do the same for sales.order_items -> mini_mart.order_items.
- d)** : Do the same for customers.customers -> mini_mart.customers.
- e)** : List ALL columns of your mini_mart schema with the introspection query below, then run it for RetailMart's 'sales' schema and compare.

Hint: Match the real COLUMN NAMES exactly (cust_id, prod_id, net_total...). Every column is TEXT in your copy - tomorrow you will ALTER them to match the real types.

Introspection Query

✓ ANSWER

```
SELECT table_name, column_name, data_type
FROM information_schema.columns
WHERE table_schema = 'mini_mart'
ORDER BY table_name, ordinal_position;
```


Answer (1/2)

✓ ANSWER

```
CREATE SCHEMA IF NOT EXISTS mini_mart;
```

```
CREATE TABLE IF NOT EXISTS mini_mart.orders (  
  order_id      TEXT,  
  cust_id       TEXT,  
  store_id      TEXT,  
  order_date     TEXT,  
  order_status  TEXT,  
  gross_total    TEXT,  
  discount_amount TEXT,  
  net_total     TEXT,  
  payment_mode_id TEXT  
);
```

```
CREATE TABLE IF NOT EXISTS mini_mart.order_items (  
  order_item_id TEXT,  
  order_id       TEXT,  
  prod_id        TEXT,  
  quantity       TEXT,  
  unit_price     TEXT,  
  gross_amount   TEXT,  
  discount_amount TEXT,
```

Answer (2/2)



ANSWER

```
net_amount      TEXT
);

CREATE TABLE IF NOT EXISTS mini_mart.customers (
  customer_id    TEXT,
  first_name     TEXT,
  last_name      TEXT,
  email          TEXT,
  phone          TEXT,
  registration_date TEXT,
  tier            TEXT,
  tier_updated_at TEXT
);

SELECT table_name, column_name, data_type
FROM information_schema.columns
WHERE table_schema = 'mini_mart'
ORDER BY table_name, ordinal_position;
```

The Empty College Analogy

You just built the college schema - students, courses, teachers, fee_payments. Beautiful structure. But there is not a single student in it. No courses listed. The system is structurally perfect and completely useless.

DML is where you START USING your tables: INSERT adds rows, UPDATE changes them, DELETE removes them, TRUNCATE empties a table. Every row in every database on Earth got there because someone ran an INSERT.

Reminder: your columns are all TEXT today, so you type EVERY value as text - including ids ('1', '2', '3'). Tomorrow, when columns get proper types and auto-IDs (SERIAL), the database fills ids for you.

INSERT INTO - Adding Data (1/2)

Technical: INSERT INTO adds one or more new rows. You name the columns, then supply matching VALUES.

Layman's: INSERT is like filling out a new row in Excel - you type the values for each column.

[SYNTAX: INSERT INTO]

-- Single row (name the columns, then the values)

```
INSERT INTO schema.table (col1, col2, col3)
```

```
VALUES ('value1', 'value2', 'value3');
```

-- Multiple rows in one statement

```
INSERT INTO schema.table (col1, col2)
```

```
VALUES ('a', '1'), ('b', '2'), ('c', '3');
```

-- INSERT ... RETURNING (see what was inserted)

```
INSERT INTO schema.table (col1, col2)
```

```
VALUES ('value1', '42')
```

```
RETURNING col1, col2;
```

INSERT INTO - Adding Data (2/2)

-- Every value is in single quotes today - all columns are TEXT.

[Worked Example - Fill college.students]

```
INSERT INTO college.students
  (student_id, first_name, last_name, email, date_of_birth, is_active)
VALUES ('1', 'Aarav', 'Sharma', 'aarav@college.edu', '2004-03-15', 'true');
```

-- Multi-row: add two more students at once

```
INSERT INTO college.students
  (student_id, first_name, last_name, email, date_of_birth, is_active)
VALUES
  ('2', 'Diya', 'Patel', 'diya@college.edu', '2003-11-02', 'true'),
  ('3', 'Kabir', 'Singh', 'kabir@college.edu', '2004-07-21', 'false');
```

-- See the row you just added

```
INSERT INTO college.students
  (student_id, first_name, last_name, email, date_of_birth, is_active)
VALUES ('4', 'Meera', 'Iyer', 'meera@college.edu', '2003-05-09', 'true')
RETURNING student_id, first_name, email;
```

UPDATE - Changing Existing Data (1/2)

Technical: UPDATE modifies existing rows. SET gives the new value(s); WHERE chooses which rows change. NO WHERE = EVERY row changes.

Layman's: UPDATE is 'Find and Replace' for your table - but if you forget WHERE, it replaces in EVERY row.

UPDATE - Changing Existing Data (2/2)

[SYNTAX & Worked Example]

-- Basic UPDATE with WHERE

```
UPDATE college.students  
SET email = 'aarav.sharma@college.edu'  
WHERE student_id = '1';
```

-- Change several columns at once

```
UPDATE college.students  
SET is_active = 'false', last_name = 'Sharma-Verma'  
WHERE student_id = '3';
```

-- RETURNING shows what changed

```
UPDATE college.students  
SET is_active = 'true'  
WHERE student_id = '3'  
RETURNING student_id, first_name, is_active;
```

Feeling the Pain of All-TEXT (and why proper data types fix it)

Because every column is TEXT today, two things feel awkward - on purpose:

- Math is clumsy. To raise a fee by 10% you cannot write `amount * 1.1` - `amount` is TEXT. You would have to cast: `amount = (amount::numeric * 1.1)::text`. Ugh.
- Comparisons sort like words, not numbers. `WHERE amount > '900'` compares TEXT, so '1000' is treated as LESS than '900' (it compares '1' vs '9'). Surprising and wrong for numbers.

That discomfort is the whole motivation for tomorrow: Data Types. You will ALTER `amount` to NUMERIC, date columns to DATE, `is_active` to BOOLEAN - and suddenly math and comparisons just work.

DELETE - Removing Data

Technical: DELETE FROM removes rows. WHERE chooses which rows. NO WHERE = ALL rows gone (the table structure stays).

Layman's: DELETE removes specific rows from a spreadsheet. The sheet and its headers remain. Forget WHERE and every row vanishes - but the empty sheet stays.

[SYNTAX & Worked Example]

-- Delete specific rows

```
DELETE FROM college.students WHERE student_id = '4';
```

-- Preview FIRST - what will this delete?

```
SELECT * FROM college.students WHERE is_active = 'false';
```

-- Happy with the preview? Delete with RETURNING to see what went

```
DELETE FROM college.students
```

```
WHERE is_active = 'false'
```

```
RETURNING student_id, first_name;
```

-- Delete ALL rows (table stays, structure intact) - dangerous!

```
-- DELETE FROM college.students;
```

TRUNCATE - The Nuclear Option

Technical: TRUNCATE removes ALL rows without scanning them one by one - far faster than DELETE for big tables. Structure stays. It is DDL, but in PostgreSQL it can be rolled back inside a transaction.

Layman's: DELETE erases each answer on an exam sheet one by one. TRUNCATE bins the sheet and hands you a fresh blank. Both give a blank sheet; TRUNCATE is instant.

[SYNTAX & Worked Example]

```
-- Empty a table fast (structure stays)
SELECT count(*) FROM college.students;    -- e.g. 3
TRUNCATE TABLE college.students;
SELECT count(*) FROM college.students;    -- 0

-- Table still works - insert fresh rows
INSERT INTO college.students (student_id, first_name)
VALUES ('1', 'Fresh Start');
```

-- Note: RESTART IDENTITY resets a SERIAL counter to 1.
-- Your columns are TEXT today (no SERIAL), so you will use that tomorrow.

DROP vs TRUNCATE vs DELETE

Q: 'What is the difference between DROP, TRUNCATE, and DELETE?'

A: 'DROP removes the table itself - structure and data gone, the table ceases to exist. TRUNCATE removes ALL rows but keeps the structure - fast, no WHERE. DELETE removes specific rows via a WHERE condition - precise, and you can roll it back inside a transaction. DROP and TRUNCATE are DDL; DELETE is DML. Memory hook: DROP = demolish the room, TRUNCATE = empty the room, DELETE = remove specific furniture.'

PRACTICE BLOCK 2 - DML: Putting Data In

Still in accio_NN, still all-TEXT columns. Build on the tables from Block 1. Do NOT touch RetailMart.

SCENARIO: Admissions are open at Prestige College. Time to put your first students and courses into the tables you built.

YOUR TASK:

a) : INSERT one student into college.students (all 6 columns) - student_id '1'.

b) : INSERT three more students in a SINGLE multi-row INSERT (ids '2','3','4').

c) : INSERT two courses into college.courses.

d) : Use RETURNING on one INSERT to see what was added.

Hint: Every value is in single quotes - all columns are TEXT. Multi-row INSERT: VALUES (...), (...), (...).

Answer



ANSWER

```
INSERT INTO college.students
(student_id, first_name, last_name, email, date_of_birth, is_active)
VALUES ('1', 'Aarav', 'Sharma', 'aarav@college.edu', '2004-03-15', 'true');
```

```
INSERT INTO college.students
(student_id, first_name, last_name, email, date_of_birth, is_active)
VALUES
('2', 'Diya', 'Patel', 'diya@college.edu', '2003-11-02', 'true'),
('3', 'Kabir', 'Singh', 'kabir@college.edu', '2004-07-21', 'false'),
('4', 'Meera', 'Iyer', 'meera@college.edu', '2003-05-09', 'true');
```

```
INSERT INTO college.courses (course_id, course_name, credits, department)
VALUES ('1', 'Database Systems', '4', 'Computer Science')
RETURNING course_id, course_name;
```

```
INSERT INTO college.courses (course_id, course_name, credits, department)
VALUES ('2', 'Data Structures', '4', 'Computer Science');
```

SCENARIO: Records change. A student updates their email, one drops out, and the registrar wants to deactivate a specific student - carefully, by id.

YOUR TASK:

- a)** : UPDATE student_id '1' - change email to 'aarav.sharma@college.edu'.
- b)** : UPDATE student_id '3' - set is_active to 'false' AND last_name to 'Singh-Rao' in ONE statement.
- c)** : Preview with SELECT, then DELETE student_id '4', using RETURNING to confirm what went.
- d)** : Verify the remaining rows with a SELECT.

Hint: Always UPDATE/DELETE by a single id to affect exactly one row. SET col1 = .., col2 = .. updates multiple columns at once.

Answer



```
UPDATE college.students  
SET email = 'aarav.sharma@college.edu'  
WHERE student_id = '1';
```

```
UPDATE college.students  
SET is_active = 'false', last_name = 'Singh-Rao'  
WHERE student_id = '3';
```

```
-- Preview first  
SELECT * FROM college.students WHERE student_id = '4';
```

```
DELETE FROM college.students  
WHERE student_id = '4'  
RETURNING student_id, first_name;
```

```
SELECT student_id, first_name, is_active FROM college.students  
ORDER BY student_id;
```


SCENARIO: The fee desk opens. You record fee payments, fix a wrongly-entered amount, remove a duplicate, then a new semester starts and the fees table must be wiped for a fresh import.

YOUR TASK:

- a)** : INSERT 4 rows into college.fee_payments (mix of tuition/hostel/library).
- b)** : UPDATE one payment whose amount was typed wrong.
- c)** : DELETE the duplicate payment (payment_id '3') with RETURNING.
- d)** : New semester: TRUNCATE college.fee_payments, then prove it is empty with count(*)

Hint: TRUNCATE clears the whole table fast and keeps the structure - perfect before a fresh bulk import. amount stays TEXT today; that is fine.

Answer



ANSWER

```
INSERT INTO college.fee_payments
(payment_id, student_id, fee_type, amount, payment_date,
 payment_mode, academic_year, semester)
VALUES
('1', '1', 'tuition', '85000', '2026-01-05', 'UPI', '2025-2026', '1'),
('2', '2', 'hostel', '45000', '2026-01-06', 'Card', '2025-2026', '1'),
('3', '2', 'hostel', '45000', '2026-01-06', 'Card', '2025-2026', '1'),
('4', '3', 'library', '2000', '2026-01-07', 'Cash', '2025-2026', '1');
```

-- (b) amount was wrong for payment 1

```
UPDATE college.fee_payments
SET amount = '90000'
WHERE payment_id = '1';
```

-- (c) row 3 is a duplicate of row 2

```
DELETE FROM college.fee_payments
WHERE payment_id = '3'
RETURNING payment_id, student_id, amount;
```

-- (d) new semester: wipe for fresh import

```
TRUNCATE TABLE college.fee_payments;
SELECT count(*) FROM college.fee_payments;    -- 0
```

SCENARIO: You prototype a tiny orders table from scratch, load it, run a full set of DML operations - and discover firsthand why all-TEXT columns make number work painful. This is the bridge to tomorrow.

YOUR TASK:

a) : Create schema shop and table shop.mini_orders (order_id, customer_name, status, total) - all TEXT.

b) : INSERT 5 orders (totals like '900', '1000', '250', '4999', '120').

c) : Mark order '1' as 'Delivered' (UPDATE).

d) : Run SELECT ... WHERE total > '900' and look at the result - notice '1000' is MISSING (TEXT compares like words!).

e) : DELETE all 'Cancelled' orders. Then TRUNCATE the table for a clean slate.

Hint: The WHERE total > '900' surprise is the lesson: TEXT sorts character-by-character, so '1000' < '900'. Tomorrow you ALTER total to NUMERIC and the comparison behaves like real numbers.

Answer (1/2)



```
CREATE SCHEMA IF NOT EXISTS shop;
```

```
CREATE TABLE IF NOT EXISTS shop.mini_orders (  
    order_id      TEXT,  
    customer_name TEXT,  
    status        TEXT,  
    total         TEXT  
);
```

```
INSERT INTO shop.mini_orders (order_id, customer_name, status, total)  
VALUES  
    ('1', 'Aarav', 'Pending', '900'),  
    ('2', 'Diya', 'Pending', '1000'),  
    ('3', 'Kabir', 'Cancelled', '250'),  
    ('4', 'Meera', 'Pending', '4999'),  
    ('5', 'Ravi', 'Cancelled', '120');
```

```
UPDATE shop.mini_orders SET status = 'Delivered' WHERE order_id = '1';
```

```
-- SURPRISE: '1000' is NOT returned - TEXT compares char-by-char,  
-- so '1000' < '900'. This is the motivation for Data Types (tomorrow).  
SELECT order_id, total FROM shop.mini_orders WHERE total > '900';
```

Answer (2/2)

✓ ANSWER

```
DELETE FROM shop.mini_orders WHERE status = 'Cancelled';  
TRUNCATE TABLE shop.mini_orders;
```

Key Takeaways (1/2)

- 1. DDL = Structure:** CREATE builds objects, DROP destroys them. CREATE DATABASE / SCHEMA / TABLE define the blueprint - what tables exist, what columns, organized into schemas.
- 2. Every column is TEXT today:** We focus on STRUCTURE, not type selection. CREATE TABLE is just 'column_name TEXT' for each column. Types come tomorrow.
- 3. Safety habits:** Always specify the schema. Use IF NOT EXISTS for idempotent scripts. DROP is permanent - practice only in accio_NN, never on RetailMart.

Key Takeaways (2/2)

4. DML = Data: INSERT adds rows (single, multi-row, RETURNING). UPDATE changes rows. DELETE removes rows. All take WHERE - and forgetting it hits EVERY row.

5. ALWAYS WHERE: UPDATE/DELETE without WHERE changes/removes every row. Preview with SELECT first, then run the change.

6. DROP vs TRUNCATE vs DELETE: DROP = demolish the room (structure gone). TRUNCATE = empty the room fast (structure stays). DELETE = remove specific furniture (WHERE).

7. The all-TEXT pain is the point: Math and comparisons are awkward on TEXT. That discomfort sets up tomorrow's Data Types lesson, where you ALTER columns to their proper types.

DDL & DML - Quick Reference

-- DDL Commands

```
CREATE DATABASE db_name;
CREATE SCHEMA schema_name;
CREATE SCHEMA IF NOT EXISTS schema_name;
CREATE TABLE schema.table (col TEXT, col TEXT, ...);
CREATE TABLE IF NOT EXISTS schema.table (...);
DROP TABLE schema.table;
DROP TABLE IF EXISTS schema.table;
DROP SCHEMA IF EXISTS schema CASCADE;
```

-- INSERT

```
INSERT INTO schema.table (c1, c2)
VALUES ('a', 'b');
-- multi-row:
INSERT INTO schema.table (c1, c2)
VALUES ('a','1'), ('b','2'), ('c','3');
-- see what was added:
INSERT INTO ... VALUES (...) RETURNING c1, c2;
```

-- UPDATE

```
UPDATE schema.table
SET col1 = 'new', col2 = 'new2'
WHERE id_col = '1';
-- ALWAYS include WHERE!
UPDATE ... SET ... WHERE ... RETURNING col;
```


DDL & DML - Quick Reference

-- DELETE / TRUNCATE

```
DELETE FROM schema.table WHERE id_col = '1';
DELETE FROM ... WHERE ... RETURNING col;
-- empty the whole table fast (structure stays):
TRUNCATE TABLE schema.table;
```

-- Inspect Your Work

```
\d schema.table          -- columns (psql)
SELECT * FROM schema.table;
SELECT count(*) FROM schema.table;
SELECT table_name, column_name, data_type
FROM information_schema.columns
WHERE table_schema = 'your_schema';
```

-- Golden Rules

```
Every column needs a TYPE (TEXT today)
Always name the schema: schema.table
IF NOT EXISTS / IF EXISTS for safe re-runs
UPDATE & DELETE: never forget WHERE
Preview with SELECT before UPDATE/DELETE
Never change RetailMart - build in accio_NN
```

INTERVIEW QUESTION BANK - DDL & DML

Practice answering OUT LOUD. Interviews test how clearly you explain, not just whether you know.

Interview Questions - DDL & DML (1/5)

Q1: 'What is DDL? Give 3 examples.'

A: 'DDL is Data Definition Language - commands that define or change database STRUCTURE. CREATE TABLE builds a table, ALTER TABLE modifies one, DROP TABLE destroys one. DDL changes the blueprint, not the data inside.'

Q2: 'What is DML? Give 3 examples.'

A: 'DML is Data Manipulation Language - commands that work with the DATA in tables. INSERT adds rows, UPDATE changes rows, DELETE removes rows. The structure stays the same; only the rows change.'

Interview Questions - DDL & DML (2/5)

Q3: 'Difference between DROP TABLE and DELETE FROM?'

A: 'DROP is DDL - it destroys the whole table, structure and data; the table no longer exists. DELETE is DML - it removes specific rows but the table stays, and you can INSERT into it again afterwards.'

Q4: 'Difference between DROP, TRUNCATE, and DELETE?'

A: 'DROP destroys the table entirely. TRUNCATE removes ALL rows but keeps the structure - fast, no WHERE. DELETE removes specific rows via WHERE. DROP and TRUNCATE are DDL; DELETE is DML.'

Interview Questions - DDL & DML (3/5)

Q5: 'What is a schema in PostgreSQL? Why not just use public?'

A: 'A schema is a namespace inside a database that groups related tables - like folders. Without schemas, every table lands in one flat public schema. With schemas like sales, customers, finance, tables are organized, and two schemas can hold tables with the same name without clashing. RetailMart uses 16 schemas.'

Q6: 'What does IF NOT EXISTS do, and why use it?'

A: 'CREATE TABLE IF NOT EXISTS creates the table only if it is not already there - otherwise it skips with a notice instead of erroring. It makes setup and migration scripts idempotent: safe to run repeatedly.'

Interview Questions - DDL & DML (4/5)

Q7: 'Why is UPDATE without WHERE dangerous?'

A: 'Without WHERE, UPDATE changes EVERY row in the table. A missing WHERE on SET price = 0 zeroes every price. Always include WHERE, and preview the rows with a SELECT using the same condition first.'

Q8: 'What does INSERT ... RETURNING do?'

A: 'RETURNING gives back columns from the rows you just inserted (or updated/deleted) in the same statement. It is great for grabbing an auto-generated id without a second query - a PostgreSQL feature MySQL lacks.'

Q9: 'Is TRUNCATE DDL or DML? Can it be rolled back?'

A: 'TRUNCATE is classified as DDL. In PostgreSQL it CAN be rolled back inside a transaction (BEGIN ... ROLLBACK), unlike some other databases where it auto-commits.'

Interview Questions - DDL & DML (5/5)

Q10: 'Walk me through building and filling a table.'

A: 'CREATE SCHEMA, then CREATE TABLE with the columns I need (DDL). Then INSERT rows (DML). I UPDATE rows to correct or change values, DELETE rows I no longer want - always with WHERE - and TRUNCATE if I need to wipe and reload.'

DDL & DML Self-Check (12 Questions)

Close your notes. Answer from memory:

- 1. What does DDL stand for? Name 3 DDL commands.
- 2. What does DML stand for? Name 3 DML commands.
- 3. Write the CREATE TABLE syntax from memory (3 TEXT columns).
- 4. Why is every column TEXT today? What changes tomorrow?
- 5. Difference between DROP TABLE, TRUNCATE, and DELETE?
- 6. What is a schema and why use one instead of public?
- 7. What does CREATE TABLE IF NOT EXISTS do?
- 8. Write a single-row INSERT and a multi-row INSERT.
- 9. Why is UPDATE without WHERE dangerous? How do you stay safe?
- 10. What does RETURNING give you?
- 11. Why does WHERE total > '900' miss the row with '1000'?
- 12. What is CASCADE in DROP SCHEMA?

Score 10+ -> ready for Data Types & Constraints. Below 10 -> re-read the weak sections.

DDL & DML Take-Home Challenge

Complete these BEFORE Data Types & Constraints:

- 1. Build 3 Tables:** : In `accio_NN`, create a mini e-commerce schema with products (7+ columns), customers (6+ columns), orders (7+ columns). ALL columns TEXT. Use `schema.table` names and IF NOT EXISTS.
- 2. Fill Them:** : INSERT at least 5 rows into each table - use a multi-row INSERT at least once, and RETURNING at least once.
- 3. Change Data:** : Write 3 UPDATE statements (each with WHERE) and 2 DELETE statements (each with WHERE). Preview each with a SELECT first.
- 4. Reset:** : TRUNCATE one table, then prove it is empty with `count(*)`.
- 5. Inspect:** : Run the `information_schema.columns` query on your schema and note that EVERY column shows `data_type = text`. Tomorrow you will fix that.

Tomorrow: Data Types (ALTER those TEXT columns to INT / NUMERIC / DATE / BOOLEAN) and Constraints & Keys (PRIMARY KEY, FOREIGN KEY, NOT NULL, UNIQUE, CHECK).

YOUR SQL JOURNEY



SQL Intro



Installation



DDL + DML <- YOU ARE HERE

← HERE



Data Types + Constraints



Filtering (first DQL!)



Joins



CTEs



Window Functions



Indexing



Views



Cohort/RFM



Capstone



WhatsApp Announcement

Tomorrow you give your columns their PROPER types and add RULES that protect your data:

- Data Types - INT, NUMERIC, DATE, TIMESTAMP, BOOLEAN, and friends
- ALTER TABLE - convert today's TEXT columns into their proper types
- NOT NULL, DEFAULT, UNIQUE, CHECK - data integrity constraints
- PRIMARY KEY - uniquely identify every row
- FOREIGN KEY - link tables together (the 'relational' part!)

Today you built the structure and filled it with data. Tomorrow you make that data CORRECT by type and SAFE by rule.

- See you on Data Types, Constraints & Keys! Siraj Sir